

# “It Was Not Your Fault” – Emotional Awareness Improves Collaborative Robots

Mahni Shayganfar  
Worcester Polytechnic Institute  
100 Institute Road  
Worcester, MA 01609  
mshayganfar@wpi.edu

Candace L. Sidner  
Worcester Polytechnic Institute  
100 Institute Road  
Worcester, MA 01609  
sidner@wpi.edu

Charles Rich  
Worcester Polytechnic Institute  
100 Institute Road  
Worcester, MA 01609  
rich@wpi.edu

Benjamin L. Hylák  
Worcester Polytechnic Institute  
100 Institute Road  
Worcester, MA 01609  
bhylak@wpi.edu

## ABSTRACT

We have conducted a user study to investigate the importance of emotional awareness and the underlying affect-driven processes during a human-robot collaboration. The goal of this user study was twofold: (1) Investigating the overall functionality of the mechanisms and the underlying algorithms in our architecture, (2) Evaluating human’s willingness and assessment of collaboration with an emotion-aware and an emotion-ignorant robot. We designed a simple table top game to simulate the collaborative environment in which a participant and the robot were “installing” a solar panel together. The result of our user study shows a significant difference between humans’ preference of working with an emotion-aware robot during collaboration.

## CCS Concepts

•Computer systems organization → Embedded systems; Redundancy; Robotics; •Networks → Network reliability;

## Keywords

Human-Robot Collaboration, Affect-Driven Processes, Emotion-Awareness

## 1. INTRODUCTION

The idea of robots or other intelligent agents living in a human environment has been a persistent dream from science fiction books to artificial intelligence and robotic laboratories. Collaborative robots are expected to become an integral part of humans’ environment to accomplish their industrial and household tasks. In these environments, humans will be involved in robots’ operations and decision-making processes. The involvement of humans influences the efficiency of robots’ interaction and performance, and

makes the robots sensitive to humans’ cognitive abilities and internal states.

We believe that collaborative robots need to take into account humans’ internal states while making decisions during collaboration. Humans express emotions to reveal their internal states in social contexts including collaboration [4]. Due to the existence of such expressions robot’s emotional-awareness can improve the quality of collaboration in terms of humans’ perception of performance and preferences. Hence, collaborative robots require to include affect-driven mechanisms in their decision making processes to be able to interpret and generate appropriate responses and behaviors. Our aim in this work was to study the importance of emotional awareness and the underlying affect-driven processes in human-robot collaboration. We examined how emotional-awareness impacts different aspects of humans’ preferences by comparing the results from our participants collaborating with an emotion-aware and an emotion-ignorant robot.

This work is implemented as part of a larger effort to assess affect-driven collaborative robots which are capable of generating and recognizing emotions in order to be better collaborators. This work is based on the development of *Affective Motivational Collaboration Theory* which is built on the foundations of the *Shared-Plans* theory of collaboration [8] and the *cognitive appraisal* theory of emotions [7].

## 2. RELATED WORK

There are many research areas, including robotics and autonomous agents, that employ the structure and/or functions of emotions in their work with a variety of motivations behind modeling emotions [28]. In [4] authors surveyed some of the principle research in social robotics and its applications in Human-Robot Interaction. We can see the application of emotion theories in designing robots capable of learning from humans [3], robots capable of expressing emotions [5] [21] and social behaviors [18], as well as robots which can convey certain types of emotion products, e.g., empathy [12]. There are also several works in which researchers have explored the human’s affective state as a mechanism to adapt the robot’s behaviors during the interaction [2] [14].

Many of the computational models of emotions and their applications are derived from appraisal theories of emotion making appraisal as the central process in their architectures [1] [16] [17] [22]. Appraisal is usually modeled as the cause of emotion being derived via simple rules on a set of appraisal variables. In robotics,

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

WOODSTOCK ‘97 El Paso, Texas USA

© 2016 ACM. ISBN 123-4567-24-567/08/06...\$15.00

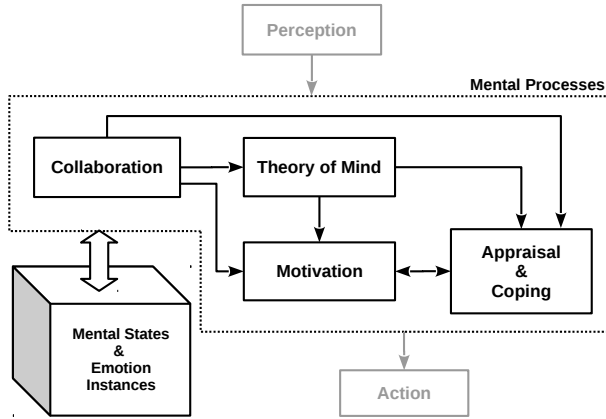
DOI: 10.475/123\_4

appraisal theory has been used to establish and maintain a better interaction between a robot and a human [6] [10] [19] [27]. There are other models of emotions that have been also used in robotics and human-robot interaction [11] [13] [29].

There are also other examples that researchers focus on the applications of emotions in human-robot collaboration. For instance, in [9] researchers use robot's emotional expression as a feedback to the human to improve the quality of collaboration. In [15] authors introduce some theoretical concepts that affective collaborative robots can enhance joint human-robot performance by adapting the robot's role and interaction to the human's affective state. This work does not provide any details of implementation and how these theoretical concepts can lead to a better human-robot collaboration. However, little effort has been put on development of functions of emotions and their applications in decision making and emotional-awareness processes of collaborative robots.

### 3. AFFECTIVE MOTIVATIONAL COLLABORATION THEORY

Affective Motivational Collaboration Theory deals with the interpretation and prediction of observable behaviors in a dyadic collaboration [24]. The theory focuses on the processes regulated by emotional states. The observable behaviors represent the outcome of reactive and deliberative processes related to the interpretation of the self's relationship to the environment. Affective Motivational Collaboration Theory aims to explain both rapid emotional reactions to events as well as slower, more deliberative responses. The reactive and deliberative processes are triggered by two types of events: *external* events, such as the other's *utterances* and *primitive actions*, and *internal* events, comprising changes in the self's mental states, such as belief formation and emotional changes. The theory explains how emotions regulate the underlying processes when these events occur. It also elucidates the role of *motives* as goal-driven emotion-regulated constructs with which a robot can form new intentions to cope with events.



**Figure 1: Computational framework based on Affective Motivational Collaboration Theory (arrows indicate primary influences between mechanisms).**

Our focus is on the mechanisms depicted as mental processes in Figure 1 along with the mental states. Each mechanism includes one or more processes in our architecture. For instance, the *Collaboration* mechanism includes processes such as *Focus Shifting* and *Constraint Management*, while the *Appraisal* mechanism includes processes to compute the values for different appraisal variables. The *mental states* includes self's (robot's) beliefs, intentions,

motives, goals and emotion instances as well as the anticipated mental states of the other (human). The *Collaboration* mechanism maintains constraints on actions, including task states and the ordering of tasks, and provides processes to update and monitor the shared plan. The *Appraisal* mechanism is responsible for evaluating changes in the self's mental states, the anticipated mental states of the other, and the state of the collaboration environment. The *Coping* mechanism provides the self with different coping strategies associated with changes in the self's mental states with respect to the state of the collaboration. The *Motivation* mechanism operates whenever the self a) requires a new motive to overcome an internal impasse in an ongoing task, or b) wants to provide an external motive to the other when the other faces a problem in a task. The *Theory of Mind* mechanism infers a model of the other's anticipated mental state. The self progressively updates this model during the collaboration.

#### 3.1 Mental States

The mental states shown in Figure 1 comprise the knowledge base required for all the mechanisms in the overall model. Mental states are conscious states of mind providing the content for cognitive processes. These mental states possess attributes, each of which provides a unique interpretation of the related cognitive entities. The self uses these attributes whenever there is an arbitration in the internal cognitive processes.

*Motives* are mental constructs which can initiate, direct and maintain goal-directed behaviors. They are created by the emotion-regulated Motivation mechanism. Motives can cause the formation of a new intention for the robot according to: a) its own emotional states, b) its own private goal, c) the collaboration (shared) goal, and d) other's anticipated beliefs. Motives possess a set of attributes. The Motivation mechanism compares motives based on the quality of these attributes and chooses the one which is the most related to the current state of the collaboration.

*Intentions* are mental constructs directed at goals and future actions. They play an essential role in taking actions according to the collaboration plan as well as behavior selection in the Coping mechanism.

*Goals* help the robot to create and update its collaboration plan according to the current private and shared goal content and structure. Goals direct the formation of intentions to take appropriate corresponding actions during collaboration.

*Emotions* in mental states are emotion instances that are elicited by the Appraisal mechanism, e.g., *Joy*, *Anger*, *Hope*, *Worry*.

### 4. COLLABORATION

The Collaboration mechanism constructs a hierarchy of goals associated with tasks in a hierarchical task network (see Figure 3), and also maintains the constraints and other required details of the collaboration including the inputs and outputs of individual tasks, the *preconditions* (specifying whether it is appropriate to perform a task), and the *postconditions* (specifying whether a just-completed task was successful). Collaboration also monitors the focus of attention, which determines the salient objects, properties and relations at each point, and shifts the focus of attention during the interaction.

### 5. APPRAISAL PROCESSES

We discuss two appraisal variables in a collaboration context, i.e., *relevance* (since other appraisals are only computed for relevant events), and *controllability* (since it is associated with the agent's coping ability). There are other appraisal variables intro-

duced in psychological [20] and computational literature [7]. We have implemented other appraisal variables such as *expectedness* [25] and *desirability* [23] which do not appear in this paper due to space limitations. The algorithms in this section use mental states of the robot (see Section 3.1) which are formed based on the collaboration structure.

## 5.1 Relevance

Relevance as an appraisal variable measures the significance of an event for the self. An event can be evaluated to be relevant if it has a non-zero utility [17]. Relevance is an important appraisal variable since the other appraisal variables are meaningful only for relevant events. However, the utility of an event during collaboration is influenced by the other collaborator’s actions and mental states, because there is a commitment between collaborators to achieve the shared goal based on the shared plan. Other appraisal models only consider the utility of an event based on the self’s (robot’s) goal and plan.

## 5.2 Controllability

Controllability is the extent to which an event can be influenced; it is associated with a robot’s ability to cope with an event [7]. Thus, a robot can determine whether an event’s outcome can be altered by actions under either of the collaborators’ control. In other words, controllability is a measure of a robot’s ability to maintain or change a particular state as a consequence of an event.

Controllability is important for the overall architecture. For instance, the robot can choose to ask or negotiate about a collaborative task which is not controllable, or form a new motive to establish an alternative goal for the current uncontrollable event. In general, other mechanisms in the architecture use the controllability output in their decision making processes; meanwhile controllability uses information from the collaboration structure, e.g., predecessors of a goal.

An important determinant of one’s emotional response is the sense of control over occurring events. This sense of subjective control is based on one’s reasoning about self’s power. For instance, the robustness of one’s plan for executing actions can increase one’s sense of power and subsequently the sense of control. In the collaboration context, we have translated the sense of control into a combination of four different factors including a) *agency* and b) *autonomy* of the robot, as well as the ratios of c) *successful predecessors*, and d) the *available inputs* of a given goal (i.e.,  $g_i$ ) in the shared plan.

In summary, the output of these two and other appraisal processes such as *desirability* [23] and *expectedness* [25] serves as critical input for the other mechanisms and processes (e.g., goal management [26]) of the Affective Motivational Collaboration Framework, shown in Figure 1. By providing adequate interpretation of events in the environment, the appraisal mechanism enables the robot to carry out proper collaborative behaviors.

## 6. EXPERIMENTAL SCENARIO

Our scenario was based on a table top turn-taking game that we designed to simulate installing a solar panel. Participants had to collaborate one-on-one with our robot to complete all the given tasks required to install the solar panel. All the tasks were simple picking up and placing collaborators’ available pegs on predefined spots on the board (see Figure 2). Each pick-and-place was associated with the robot’s or the participant’s task. The robot and the participants had their own unique primitive tasks that they had to accomplish in their own turn. The final goal of installing a solar panel required the robot and the participants to accomplish their




Figure 2: The layout of the available spots for the human and the robot to place their pegs during the collaboration.

own individual tasks. Failure of any task could create an impasse during the collaboration.

### 6.1 The Robot

We conducted our experiment based on a KUKA Youbot (see Figure 4). The robot was stationary on top of a desk and was able to pick up and place available pegs corresponding to the robot’s task. The robot was operated based on ROS and was receiving commands through the ROS-bridge from either our Affective Motivational Collaboration framework or Disco, see Figure XYZ.

### 6.2 Interaction Paradigms

The robot interacted via a) speech, b) the corresponding utterance on the screen, c) negative, positive and neutral expression of emotion through an emoticon on the screen. The robot used neutral expression in case of the emotion-ignorance. The interaction was controlled autonomously by the AMC framework in case of the emotion-awareness, and Disco in case of the emotion-ignorance (see Section XYZ). The reasoning about which task should be done and controlling the robot was entirely autonomous. Only the perception of the task failure or achievement by the robot or by the participant was done by a wizard monitoring the collaboration outside of the test area. The interaction was structured based on the exact same goals in an HTN for both conditions. The robot was using the same utterances in both conditions. In emotion-aware condition only if the participant was expressing a negative emotion in case of a failure the robot was using a different utterance in compare to the participant’s positive or neutral expression of emotion; i.e., the robot’s utterances were identical in emotion-aware and emotion-ignorance cases if in the latter the participant reported (expressed) a positive or a neutral emotion. At the beginning of each collaboration the robot asked each participant to achieve the overall shared goal, i.e., “installing the solar panel”. Then, before achieving a new goal, the robot informed the participant about the higher level non-primitive goal of which the primitives were going to be achieved. The same procedure was used by the robot if there was a decision for switching to another nonprimitive due to the failure of the task achieving the current goal. For example, ... (see Figure 3). After achieving a new primitive goal, the robot either informed the human keeping the ground for the next goal to achieve, or informed and passed the ground to the human to execute the next task with respect to the human’s goal. In case of the human’s turn, the robot waited for the human to do a task, then the wizard let the robot know whether the human’s goal was achieved or not. Afterwards the robot was making a decision about which goal to pursue and




Figure 3: Collaboration structure used as the task model.




Figure 4: Experimental setup.

was informing human accordingly. The same entire procedure was applied to both conditions.

### 6.3 Environment and Tasks

The environment was set up in Human-Robot Interaction lab, and included the robot, the collaboration board on top of a desk, and the participant standing in front of the robot on the other side of the board (see Figure 4). One of the experimenters monitored the interactions using a live stream of a camera in a different room. The experimenter provided only the required perception, i.e., decision on success or failure of the tasks for the robot, through the entire time of the collaboration (see Section 6.2).

The tasks were defined based on the HTN structure shown in Figure XYZ and were executed in a turn-taking fashion by either of the collaborators. For each task either the robot or the participant was responsible to pick up one of the corresponding pegs from their own inventory and place it on the right spot which was colored and tagged same as the associated peg.

## 7. EVALUATION AND RESULTS

### 7.1 Hypothesis

The non/social functions of emotions impact a collaboration process. Human collaborators prefer to collaborate with others whose

behaviors are influenced by these functions of emotions depending on the context. We developed seven hypotheses on positive influence of emotion-awareness and usefulness of emotion function during collaboration:

**Hypothesis 1.** Subjects will prefer to collaborate with the robot which is controlled by Affective Motivational Collaboration framework more than the robot which is controlled by Disco.

**Hypothesis 2.** Subjects will find a mutual understanding in the collaboration with the robot which is controlled by the AMC framework more than the robot controlled by Disco.

**Hypothesis 3.** Subjects will find working with the robot which is controlled by AMC framework less confusing than the robot controlled by Disco.

**Hypothesis 4.** Subjects will find the robot which is controlled by AMC framework understanding their goals more than the robot controlled by Disco.

**Hypothesis 5.** Subjects will find the robot which is controlled by AMC framework understanding their feelings more than the robot controlled by Disco.

**Hypothesis 6.** Subjects will find the robot which is controlled by AMC framework more helpful in compare to the robot controlled by Disco.

**Hypothesis 7.** Subjects will find the robot which is controlled by AMC framework more trustworthy than the robot controlled by Disco.

### 7.2 Procedure

Participants were first given a brief description of the purpose of the experiment. After the short introduction, they were asked to review and sign a consent form. Participants were then provided with a written instruction of their task and the rules for collaborating with the robot. Then, one of the experimenters lead them into the experiment room and asked the participants to answer pre-experiment questionnaires. Afterwards, the experimenter went through all the important details of the instructions with the participants standing in front of the collaboration board and the robot. The experimenter confirmed participants' correct understanding of the tasks and informed them with type of task failures that might occur during the collaboration. Participants were told that researchers were developing a collaborative robot and would like their help in evaluating their design. Participants were provided with identical instructions and randomly assigned to the conditions in the

experiment. They were told that, after their collaboration with the robot, they would be asked to answer a questionnaire on their experience. After completing the first round of collaboration, participants answered a post-experiment questionnaire that measured their perceptions of the robot, the task, and the collaboration procedure. After answering the first post-experiment questionnaire, participants were told that they were going to collaborate with the robot one more time and the robot might not necessarily have the same collaborative behavior. After completing the second round of collaboration, participants were asked to answer the second post-experiment questionnaire which included the same questions as the first post-experiment questionnaire. After all, participants were asked to answer an open-ended questionnaire which measured their perception of difference between two runs, their preference of collaborative robot between two runs, and their reasons of preference.

### 7.3 Measurements

In our study two basic conditions of the robot were tested: a) controlling the robot using Disco, b) controlling the robot using AMC framework. The collaborative results were measured using objective and subjective measurements.

**Objective –** We measured participants’ recall of the collaborative behaviors presented by the robot using an open-ended post-experiment questionnaire. We also specifically asked the participants what behavior of the robot did they like during their collaboration.

**Subjective –** We evaluated participants’ levels of satisfaction, trust, confusion, goal achievement, as well as mutual understanding of goals, mutual understanding of feelings, mutual agreement, and also participants’ beliefs about the efficiency of collaboration and their feeling of robot’s collaborative behaviors. Seven-point Likert scales were used in all questionnaire items.

### 7.4 Participants

A total of 37 participants participated in the experiment in 74 trials. Participants were recruited from Worcester Polytechnic Institute’s students and staffs as well as other civilians recruited from outside of the campus. The ages of the participants varied between 19 and 74 with an average of 34.2 years before our screening of 4 subjects based on our sanity check questions. After this screening the ages of the participants varied between 19 and 54 with an average of 30.8 years old. Of all the 33 participants, 21 were female and 12 were male. Each participant participated in 2 trials. In one trial the robot was controlled using Disco and in the second trial the robot was controlled using AMC framework. The order of these two trials were randomly assigned to each participant. In general we used Disco first in 16 experiments, and AMC framework first in 17 experiments.

### 7.5 Evaluation Results

**Subjective –** In our first hypothesis we predicted that the participants would prefer to collaborate with the robot controlled by AMC framework than the robot controlled by Disco. This prediction was supported by our analysis. There was a significant difference across two conditions in participants’ preference to collaborate with the robot. In general ??? out of 33 participants preferred to collaborate with the robot controlled by AMC framework; ??? out of 33 participants found at least one behavior generated by AMC framework useful for more complex tasks; ??? out of 33 participants found the robot controlled by AMC framework being able to exhibit behaviors that can prevent human errors; ??? out of 33 participants believe the robot controlled by AMC framework can exhibit behaviors that could improve efficiency of collaboration; and, ??? out of

33 participants mentioned at least one of the collaborative behaviors of the robot controlled by AMC framework as an interesting behavior during collaboration.

**Objective –**

## 8. ANALYSIS AND DISCUSSION

### 8.1 Generalization and Limitations

## 9. CONCLUSIONS

## 10. ACKNOWLEDGMENTS

## 11. REFERENCES

- [1] C. Adam and E. Lorini. A BDI emotional reasoning engine for an artificial companion. In *Workshop on Agents and multi-agent Systems for AAL and e-HEALTH (PAAMS)*, volume 430, pages 66–78. Springer, 2014.
- [2] C. Breazeal. *Designing Sociable Robots*. MIT Press, 2002.
- [3] C. Breazeal. Role of expressive behaviour for robots that learn from people. *Philosophical transactions of the Royal Society of London. Series B, Biological sciences*, 364(1535):3527–38, 2009.
- [4] C. Breazeal, A. Takanishi, and T. Kobayashi. *Social Robots that Interact with People*, pages 1349–1369. Springer Berlin Heidelberg, Berlin, Heidelberg, 2008.
- [5] D. Cameron, S. Fernando, E. Collins, A. Millings, R. Moore, A. Sharkey, V. Evers, and T. Prescott. Presence of life-like robot expressions influences children’s enjoyment of human-robot interactions in the field. In *4th International symposium on New Frontiers in Human-Robot Interaction.*, Canterbury, UK, 2015. M. Salem, A. Weiss, P. Baxter, & K. Dautenhahn (Eds.).
- [6] A. C. Gonzalez, M. Malfaz, and M. A. Salichs. An autonomous social robot in fear. *IEEE Transactions Autonomous Mental Development*, 5(2):135–151, 2013.
- [7] J. Gratch and S. C. Marsella. A domain-independent framework for modeling emotion. *Cognitive Systems Research*, 5(4):269–306, 2004.
- [8] B. J. Grosz and C. L. Sidner. Plans for discourse. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*, pages 417–444. MIT Press, Cambridge, MA, 1990.
- [9] A. I. Guha and S. Tellex. Towards Meaningful Human-Robot Collaboration on Object Placement. In *RSS Workshop on Planning for Human-Robot Interaction: Shared Autonomy and Collaborative Robotics*, 2016.
- [10] H.-R. Kim and D.-S. Kwon. Computational model of emotion generation for human-robot interaction based on the cognitive appraisal theory. *Journal of Intelligent and Robotic Systems*, 60(2):263–283, 2010.
- [11] M. Klug and A. Zell. Emotion-based human-robot-interaction. In *IEEE 9th International Conference on Computational Cybernetics (ICCC)*, 2013.
- [12] I. Leite, A. Pereira, S. Mascarenhas, C. Martinho, R. Prada, and A. Paiva. The influence of empathy in human-robot relations. *International Journal of Human-Computer Studies*, 71(3):250–260, 2013.
- [13] A. Lim and H. G. Okuno. The mei robot: Towards using motherese to develop multimodal emotional intelligence.

*IEEE Transactions Autonomous Mental Development*,  
6(2):126–138, 2014.

- [14] C. Liu and N. Sarkar. Online affect detection and robot behavior adaptation for intervention of children with autism. *IEEE TRANSACTIONS ON ROBOTICS*, 24(4):883–896, 2008.
- [15] R. Looije, M. Neerinx, and G.-J. Kruijff. Affective collaborative robots for safety crisis management in the field. In *Proceedings of the 4th International Conference on Information Systems for Crisis Response and Management (ISCRAM 2007)*. ISCRAM, 5 2007.
- [16] R. P. Marinier III, J. E. Laird, and R. L. Lewis. A computational unification of cognitive behavior and emotion. *Cognitive System Research*, 10(1):48–69, March 2009.
- [17] S. C. Marsella and J. Gratch. EMA: A process model of appraisal dynamics. *Cognitive Systems Research*, 10(1):70–90, March 2009.
- [18] A. Paiva, I. Leite, and T. Ribeiro. Emotion modeling for sociable robots. In J. G. A. K. Rafael A. Calvo, Sidney D’Mello, editor, *Handbook of Affective Computing*, pages 296–308. Oxford University Press, 2014.
- [19] M. Pontier and J. F. Hoorn. How women think robots perceive them - as if robots were men. In *International Conference on Agents and Artificial Intelligence (ICAART-2)*, pages 496–504, 2013.
- [20] K. R. Scherer, A. Schorr, and T. Johnstone. *Appraisal Processes in Emotion: Theory, Methods, Research*. Oxford University Press, 2001.
- [21] M. Shayganfar, C. Rich, and C. L. Sidner. A design methodology for expressing emotion on robot faces. In *IROS*, pages 4577–4583. IEEE, 2012.
- [22] M. Si, S. C. Marsella, and D. V. Pynadath. Modeling appraisal in theory of mind reasoning. *Autonomous Agents and Multi-Agent Systems*, 20(1):14–31, 2010.
- [23] Suppressed for Anonymity. 2015.
- [24] Suppressed for Anonymity. 2016.
- [25] Suppressed for Anonymity. 2016.
- [26] Suppressed for Anonymity. 2016.
- [27] D. Vogiatzis, C. Spyropoulos, V. Karkaletsis, Z. Kasap, C. Matheson, and O. Deroo. An affective robot guide to museums. In *Proceedings of the 4th International Workshop on Human-Computer Conversation*, 2008.
- [28] T. Wehrle. Motivations behind modeling emotional agents: Whose emotion does your robot have?, 1998.
- [29] T. Zhang, D. B. Kaber, B. Zhu, M. Swangnetr, P. Mosaly, and L. Hodge. Service robot feature design effects on user perceptions and emotional responses. *Intelligent Service Robotics*, 3(2):73–88, 2010.