

“It Was Not Your Fault” – Emotional Awareness Improves Collaborative Robots

ABSTRACT

We have conducted a user study to investigate the importance of emotional awareness and the underlying affect-driven processes during a human-robot collaboration. The goal of this user study was twofold: (1) Investigating the overall functionality of the mechanisms and the underlying algorithms in our architecture, (2) Evaluating human’s willingness and assessment of collaboration with an emotion-aware and an emotion-ignorant robot. We designed a simple table top activity to simulate the collaborative environment in which a participant and the robot were “installing” a solar panel together. The result of our user study shows that humans’ significantly prefer working with an emotion-aware robot during collaboration.

Keywords

Human-Robot Collaboration, Affect-Driven Processes, Emotion-Awareness

1. INTRODUCTION

The idea of robots or other intelligent agents living in a human environment has been a persistent dream from science fiction books to artificial intelligence and robotic laboratories. Collaborative robots are expected to become an integral part of humans’ environment to accomplish their industrial and household tasks. In these environments, humans will be involved in robots’ operations and decision-making processes. The involvement of humans influences the efficiency of robots’ interaction and performance, and makes the robots sensitive to humans’ cognitive abilities and internal states.

We believe that collaborative robots need to take into account humans’ internal states while making decisions during collaboration. Humans express emotions to reveal their internal states in social contexts including collaboration [4]. Due to the existence of such expressions robots’ emotional-awareness can improve the quality of collaboration in terms of humans’ perception of performance and preferences. Hence, collaborative robots need to include affect-driven mechanisms in their decision making processes to be able to interpret and generate appropriate responses and behaviors.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

WOODSTOCK ’97 El Paso, Texas USA

© 2016 ACM. ISBN 123-4567-24-567/08/06...\$15.00

DOI: 10.475/123_4

Our aim in this work was to study the importance of emotional awareness and the underlying affect-driven processes in human-robot collaboration. We examined how emotional-awareness impacts different aspects of humans’ preferences by comparing the results from our participants collaborating with an emotion-aware and an emotion-ignorant robot.

This work is implemented as part of a larger effort to assess affect-driven collaborative robots which are capable of generating and recognizing emotions in order to be better collaborators. This work is based on the development of [Title Suppressed For Anonymity] Theory which is built on the foundations of the *Shared-Plans* theory of collaboration [8] and the *cognitive appraisal* theory of emotions [7].

2. RELATED WORK

There are many research areas, including robotics and autonomous agents, that employ the structure and/or functions of emotions in their work with a variety of motivations behind modeling emotions [27]. In [4] authors surveyed some of the principle research in social robotics and its applications in Human-Robot Interaction. We can see the application of emotion theories in designing robots capable of learning from humans [3], robots capable of expressing emotions [5] [20] and social behaviors [18], as well as robots which can convey certain types of emotion products, e.g., empathy [12]. There are also several works in which researchers have explored the human’s affective state as a mechanism to adapt the robot’s behaviors during the interaction [2] [14].

Many of the computational models of emotions and their applications are derived from appraisal theories of emotion making appraisal the central process in their architectures [1] [16] [17] [21]. Appraisal is usually modeled as the cause of emotion, being derived via simple rules on a set of appraisal variables. In robotics, appraisal theory has been used to establish and maintain a better interaction between a robot and a human [6] [10] [19] [26]. There are other models of emotions that have also been used in robotics and human-robot interaction [11] [13] [28].

There are also other examples in which researchers focus on the applications of emotions in human-robot collaboration. For instance, in [9] researchers use a robot’s emotional expression as a feedback to the human to improve the quality of collaboration. In [15] the authors introduce some theoretical concepts including that affective collaborative robots can enhance joint human-robot performance by adapting the robot’s role and interaction to the human’s affective state. This work does not provide any details of implementation or how these theoretical concepts can lead to a better human-robot collaboration. Little effort has been put on development of functions of emotions and their applications in decision making and emotional-awareness processes of collaborative robots.

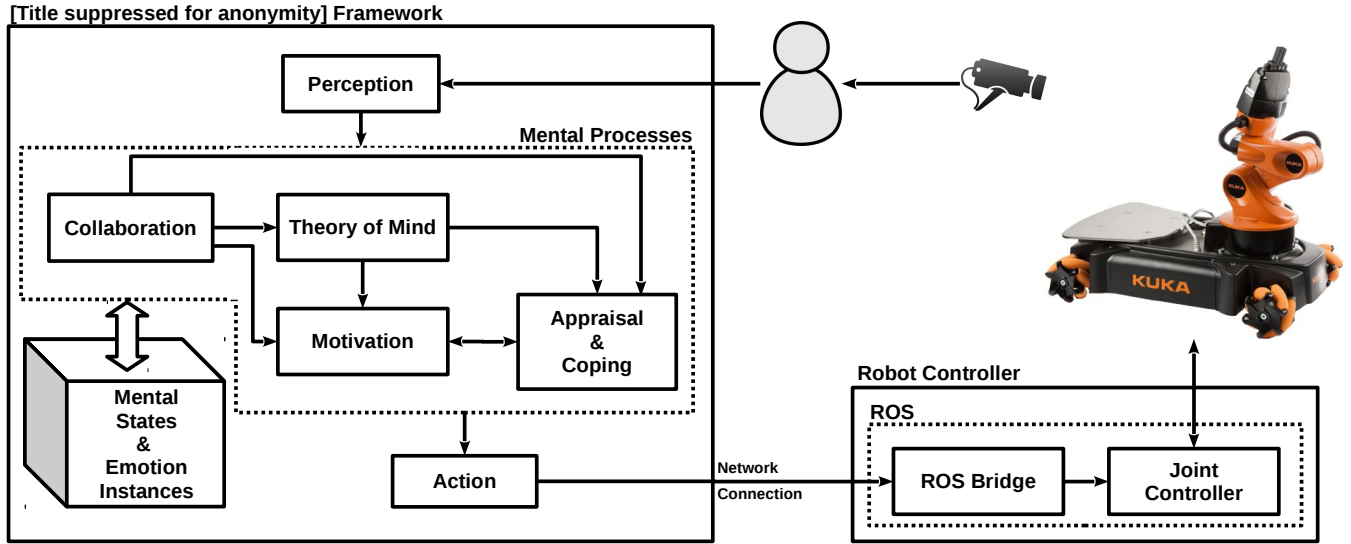


Figure 1: Computational framework based on [title suppressed for anonymity] theory (arrows indicate primary influences between mechanisms and data flow).

3. IMPLEMENTATION

The implementation of this user-study included two separated parts. The first part incorporated the [Title Suppressed For Anonymity] Framework consisting of all Mental Processes (see left-side of Figure 1) briefly described in Section 3.1. The second part was implemented to receive action commands from the framework and forward them to the robot to control joints and actuators (see right-side of Figure 1).

3.1 [Title Suppressed For Anonymity] Framework

This framework is built based on [Title Suppressed For Anonymity] Theory which deals with the interpretation and prediction of observable behaviors in a dyadic collaboration [23]. The theory focuses on the processes regulated by emotional states [24]. The observable behaviors represent the outcome of reactive and deliberative processes related to the interpretation of the self’s relationship to the environment. [Title Suppressed For Anonymity] Theory aims to explain both rapid emotional reactions to events as well as slower, more deliberative responses. The reactive and deliberative processes are triggered by two types of events: *external* events, such as the other’s *utterances* and *primitive actions*, and *internal* events, comprising changes in the self’s mental states, such as belief formation and emotional changes. The theory explains how emotions regulate the underlying processes when these events occur. It also elucidates the role of *motives* as goal-driven emotion-regulated constructs with which a robot can form new intentions to cope with events.

The framework includes the mechanisms depicted as mental processes in Figure 1 along with the mental states. The mental states shown in Figure 1 comprise the knowledge base required for all of the mechanisms in the overall model. These *mental states* include self’s (robot’s) beliefs, intentions, motives, goals and emotion instances as well as the anticipated mental states of the other (human). The details about these mental states are beyond the scope of this paper.

Each mechanism includes one or more processes in our architecture. For instance, the *Collaboration* mechanism includes pro-

cesses such as *Focus Shifting* and *Constraint Management*, while the *Appraisal* mechanism includes processes to compute the values for individual appraisal variables [22] [25]. The *Collaboration* mechanism maintains constraints on actions, including task states and the ordering of tasks, and provides processes to update and monitor the shared plan. In this user-study, the *Collaboration* mechanism uses a hierarchy of goals associated with tasks in a Hierarchical Task Network (HTN) structure depicted in Figure 2. The *Appraisal* mechanism is responsible for evaluating changes in the self’s mental states, the anticipated mental states of the other, and the state of the collaboration environment. The *Coping* mechanism provides the self with different coping strategies associated with changes in the self’s mental states with respect to the state of the collaboration. The *Motivation* mechanism operates whenever the self a) requires a new motive to overcome an internal impasse in an ongoing task, or b) wants to provide an external motive to the other (i.e. human) when the other faces a problem in a task. The *Theory of Mind* mechanism infers a model of the other’s anticipated mental state. The self progressively updates this model during the collaboration.

3.2 Robot Controller

The robot controller is comprised of two major components: 1) ROS-bridge and 2) joint controller (see Figure 1). ROS-bridge¹ provides an API to ROS functionality for non-ROS programs which enables us to send action commands from our framework (implemented in JAVA) to the robot’s joint controller. The joint controller receives action commands and translates them into actual joint and actuator commands and sends them to the robot.

4. EXPERIMENTAL SCENARIO

Our scenario was based on a table top turn-taking game that we designed to simulate the installation of a solar panel. Participants had to collaborate one-on-one with our robot to complete all the given tasks required to install the solar panel. All the tasks consisted of picking up and placing collaborators’ available pegs on predefined spots on the board (see Figure 3). Each pick-and-place

¹http://wiki.ros.org/rosbridge_suite

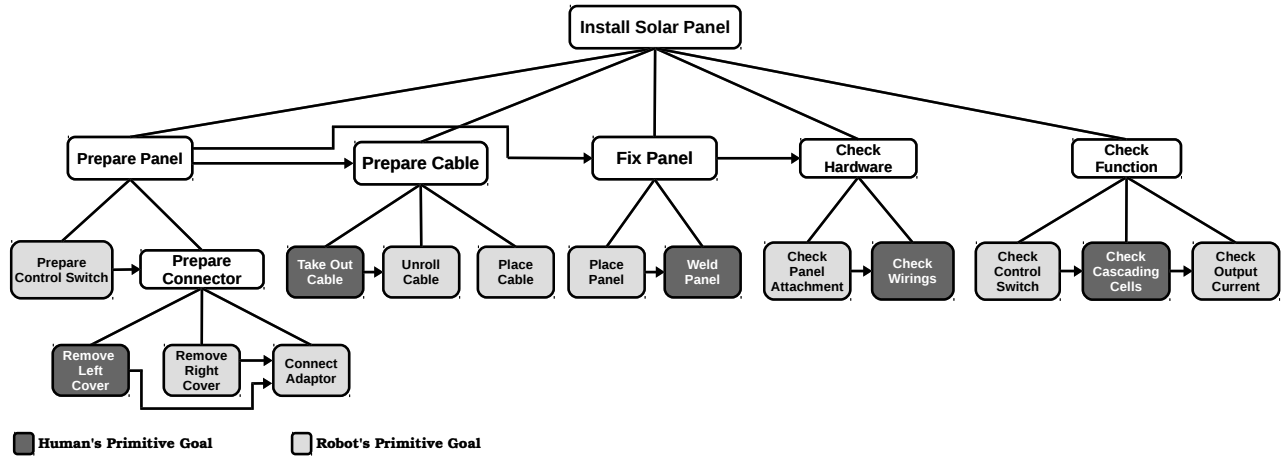


Figure 2: Collaboration structure used as the task model.

was associated with the robot’s or the participant’s task. The robot and the participants had their own unique primitive tasks that they had to accomplish in their own turn. The final goal of installing a solar panel required the robot and the participants to accomplish their own individual tasks. Failure of any task could create an impasse during the collaboration.

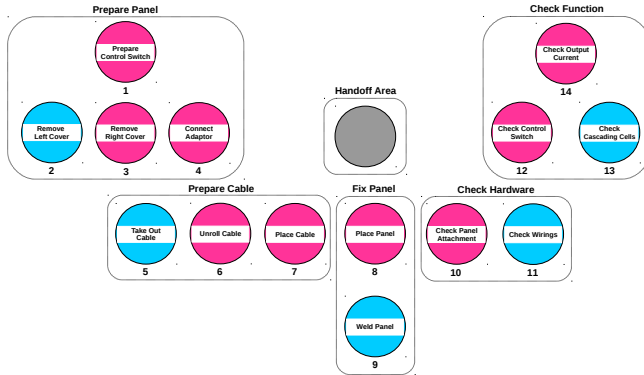


Figure 3: The layout of the available spots for the human and the robot to place their pegs during the collaboration.

4.1 The Robot

We conducted our experiment based on a KUKA Youbot (see Figure 4). The robot was stationary on top of a desk and was able to pick up and place available pegs corresponding to the robot’s task. The robot was operated based on Robot Operating System (ROS – indigo) and was receiving commands through the ROS-bridge from our [Title Suppressed For Anonymity] framework (see Figure 1). We used a touch-screen monitor (see Figure 4) to a) express robot’s positive, negative or neutral emotion through an emoticon, b) display robot’s utterances, c) control turn-taking process of the collaboration, and d) let the participants express (report) their positive, negative or neutral emotion for each turn. The robot used MaryTTS an open-source, multilingual Text-to-Speech Synthesis platform to provide corresponding speech for its utterances in English.

4.2 Interaction Paradigms

At the beginning of each collaboration the robot asked each participant to achieve the overall shared goal, i.e., “installing the solar panel”. Then, before working towards a new goal, the robot informed the participant about the higher level nonprimitive goal (e.g. Prepare Panel – see Figure ??) of which the primitives were going to be working towards. The same procedure was used by the robot if there was a decision to switch to another nonprimitive due to the failure of a task in achieving the current goal. After achieving a new primitive goal, the robot either informed the human that it would pursue the next goal, or it informed and passed the turn to the human to execute the next task with respect to the human’s goal. In case of the human’s turn, the robot waited for the human to do a task, then the wizard let the robot know whether the human’s goal was achieved or not. Afterwards the robot made a decision about which goal to pursue and informed the human accordingly. The same procedure was applied to both conditions.

The robot interacted via a) speech, b) the corresponding utterance on the screen, c) negative, positive and neutral expression of emotion through an emoticon on the screen. There were two conditions of the robot: 1) emotion-aware and 2) emotion ignorant. The robot used neutral expression in the case of emotion-ignorance. The interaction was controlled autonomously by the framework we discussed in Section 3.1 in both the emotion-ignorant and the emotion-aware cases. The reasoning about which task should be done and controlling the robot was entirely autonomous. Only the perception of the task failure or achievement by the robot or by the participant was done by a wizard monitoring the collaboration outside of the test area. The interaction was structured based on the exact same goals in an HTN for both conditions. The robot was using the same utterances in both conditions. In the emotion-aware condition the robot used a different behavior in comparison with the emotion-ignorant condition only if the participant was expressing a negative emotion in case of a failure; i.e., the robot’s utterances were identical in emotion-aware and emotion-ignorant cases if in the latter the participant reported (expressed) a positive or a neutral emotion.

Three different behaviors could be generated only in the emotion-aware condition. These three behaviors were 1) mitigating the human’s negative emotion and postponing its own task to help the human, 2) goal-management to switch to another goal which has lower cost with respect to the human’s negative emotion, and 3) task delegation to the human to overcome the impasse. In each

run, the human had two pre-coordinated task failures, and the robot had one. If the human expressed negative emotion after the first human task-failure, the robot responded by mitigating the human's negative emotion by saying "It was not your fault. I can help you with this task" and helping the human by providing a peg to fulfill the human's task. If the human expressed negative emotion after the second human task-failure, the robot informed the human that they could proceed with another task to save time while simultaneously requesting a new peg (i.e. help) from the supervisor. If the human expressed negative emotion as a result of the robot's task failure, the robot requested help from the human (who had the correct peg). In the event that the human expressed positive or neutral emotion during these three failures, the robot behaved identically in the emotion-ignorant and the emotion-aware cases, by asking the supervisor for help.

4.3 Environment and Tasks

The environment was set up in the Human-Robot Interaction lab, and included the robot, the collaboration board on top of a desk, and the participant standing in front of the robot on the other side of the board (see Figure 4). One of the experimenters monitored the interactions using a live stream of a camera in a different room. The experimenter provided only the required perception, i.e., decision on success or failure of the tasks for the robot, through the entire time of the collaboration (see Section 4.2).

The tasks were defined based on the HTN structure shown in Figure ?? and were executed in a turn-taking fashion by either of the collaborators. For each task either the robot or the participant was responsible to pick up one of the corresponding pegs from their own inventory and place it on the right spot which was colored and tagged same as the associated peg. Some pegs and corresponding spots on the board had hidden magnets which prevented the pegs from standing upright. Any peg that fell over was considered a failed task.

5. EVALUATION

5.1 Hypothesis

The non/social functions of emotions impact a collaboration process. Human collaborators prefer to collaborate with others whose behaviors are influenced by these functions of emotions depending on the context. We developed seven hypotheses on positive influence of emotion-awareness and usefulness of emotion function during collaboration:

Hypothesis 1. Subjects will feel closer (likability) to the emotion-aware robot rather than the emotion-ignorant robot.

Hypothesis 2. Subjects will find the emotion-aware robot to be more trustworthy than the emotion-ignorant robot.

Hypothesis 3. Subjects will find the emotion-aware robot to have better performance in collaboration than the emotion-ignorant robot.

Hypothesis 4. Subjects will find the emotion-aware robot to be more understanding of their feelings than the emotion-ignorant robot.

Hypothesis 5. Subjects will find the emotion-aware robot to be more understanding of their goals than the emotion-ignorant robot.

Hypothesis 6. Subjects will feel more satisfied about the collaboration when working with the emotion-aware robot rather than emotion-ignorant robot.

Hypothesis 7. Subjects will perceive higher level of mutual satisfaction with emotion-aware robot than emotion-ignorant robot.




Figure 4: Experimental setup.

5.2 Procedure

Participants were first given a brief description of the purpose of the experiment. After the short introduction, they were asked to review and sign a consent form. Participants were then provided with a written instruction of their task and the rules for collaborating with the robot. Then, one of the experimenters lead them into the experiment room and asked the participants to asked to answer pre-experiment questionnaires. Afterwards, the experimenter went through all the details of the instructions with the participants standing in front of the collaboration board and the robot. The experimenter confirmed participants' correct understanding of the tasks and informed them of types of task failures that might occur during the collaboration. Participants were told that researchers were developing a collaborative robot and would like their help in evaluating their design. Participants were provided with identical instructions and randomly assigned to the conditions in the experiment. They were told that, after their collaboration with the robot, they would be asked to answer a questionnaire on their experience. After completing the first round of collaboration, participants answered a post-experiment questionnaire that measured their perceptions of the robot, the task, and the collaboration procedure. After answering the first post-experiment questionnaire, participants were told that they were going to collaborate with the robot one more time and the robot might not necessarily have the same collaborative behavior. After completing the second round of collaboration, participants were asked to answer the second post-experiment questionnaire which consisted of the same questions as the first post-experiment questionnaire. After all, participants were asked to answer an open-ended questionnaire which measured their perception of difference between two runs, their preference of collaborative robot between two runs, and their reasons of preference.

5.3 Measurements

In our study two basic conditions of the robot were tested: a) the emotion-ignorant condition, b) the emotion-aware condition. We measured participants' recall of the collaborative behaviors presented by the robot using an open-ended post-experiment questionnaire. We also specifically asked the participants what behavior of the robot they liked during their collaboration. We also evaluated participants' levels of satisfaction, trust, confusion, goal achievement, mutual understanding of goals, mutual understanding of feelings, mutual agreement, and also participants' beliefs about the efficiency of collaboration and their feeling of robot's collaborative behaviors. Seven-point Likert scales were used in these questionnaire items.

5.4 Participants

A total of 37 participants participated in the experiment in 74 trials. Participants were recruited from Worcester Polytechnic Institute's students and staffs as well as other civilians recruited from outside of the campus. The ages of the participants varied between 19 and 74 with an average of 34.2 years before our screening of 4 subjects based on our sanity check questions. After this screening the ages of the participants varied between 19 and 54 with an average of 30.8 years old. Of the 33 participants, 21 were female and 12 were male. Each participant participated in 2 trials. In one trial the robot was aware of human's emotion and in the second trial the robot was ignoring human's emotion. The order of these two trials were randomly assigned to each participant. In general we used emotion-ignorant robot first in 16 experiments, and emotion-aware robot first in 17 experiments.

6. RESULTS

As discussed in Section 5.3, results of the user study were gathered through a 31-question Likert-scale survey that was given to each participant after each run with the robot, and through a 5-question open-ended summary questionnaire at the end of the experiment.

6.1 7-Point Likert Scale Survey Results

As mentioned previously, the 7-point Likert scale survey was administered at the end of the emotion-ignorant run and at the end of the emotion-aware run for each participant. The 31 questions are generally categorized to evaluate the humans' perceptions of the following seven categories, with 3-6 questions per group: (1) the likability of the robot (2) the level of trust the human feels in the robot (3) the human's perception of the robot's performance (4) the human's perception of the robot's understanding of human's emotions (5) the human's perception of the robot's understanding of human's and collaboration's goals and objectives (6) the human's feeling about the collaboration and (7) the human's perception of the human's and robot's mutual satisfaction with each other as collaborative partners. Of the 31 questions, 2 were chosen from each group, for a total of 14 questions, for presentation in this paper. The remaining questions were omitted due to space constraints. The questions presented are provided in Figure 6. These questions are chosen as representative of their respective groups of questions, and their results do not necessarily represent the highest levels of statistical significance.

The results were analyzed using a two-tailed paired t-test to analyze the difference of means between the emotion ignorant and the emotion-aware condition. Refer to Figure 7 for the results. As mentioned in Section 5.4, participants were randomly assigned to complete either the emotion-ignorant or the emotion-aware run first; analysis of the results revealed no statistically significant difference or consistent pattern based on which run the participant completed first.

6.1.1 Likability of the Robot

Questions 1 and 2 addressed the likability of the robot. As shown in Figure 7, participants would like to continue working with the emotion-aware robot significantly more than the emotion-ignorant robot by an average of about 1.5 points; additionally participants felt more close to the emotion-aware robot than the emotion-ignorant robot by about 2.1 points, on the 7-point scale. This supports Hypothesis 1, which stated that humans would prefer to work with the emotion-aware robot over the emotion-ignorant robot.

6.1.2 Human Trust in the Robot

Questions 3 and 4 were designed to measure the degree of trust that the human participants felt in the robot. As shown in Figure 7, participants trusted the emotion-aware robot an average of 1.5 points more than the emotion-ignorant robot, both in general and in terms of collaboration performance. In fact, in Question 4, participants rated their trust in the emotion-aware robot to perform appropriately during collaboration an average of 5.9 on a 7-point Likert scale, where 7.0 would indicate maximum trust; this indicates that the participants felt that the emotion-aware robot's collaborative performance was acceptable. This result supports Hypothesis 2, that posits that human participants would find the emotion-aware robot to be more trustworthy than the emotion-ignorant robot.

6.1.3 Perception of the Robot's Performance

Question 5 (which is reverse-scored) measures the participant's perception of repetitiveness in the robot during the collaboration. In both conditions, participants rated the robot as moderately repetitive, with the emotion-ignorant robot's average response being about 1.1 points higher than the emotion-aware. This result correlates with several of the open-ended responses which described the emotion-aware robot's behaviors as "cute" and "interesting", refer to Section 6.2. The participants also felt that the emotion-aware robot's decisions during collaboration improved their own performance, with an average rating of 5.4, while the emotion-ignorant robot only received an average rating of 3.3, indicating that participants felt it was not able to interact in such a way as to increase the human's performance; refer to results from Question 6. These results support Hypothesis 3, which posited that humans will perceive the emotion-aware robot as being more capable than the emotion-ignorant robot.

6.1.4 Robot's Understanding of Human Emotions

For Questions 7 and 8, participants ranked the emotion-aware robot's understanding of emotions 2.2 and 1.8 points higher, respectively, than the emotion-ignorant robot's understanding of emotions, supporting Hypothesis 4. This category showed the highest total difference between the emotion-ignorant and the emotion-aware robot.

6.1.5 Robot's Understanding of Human and Collaboration Goals

Question 9 was a measure of whether the human perceived that the robot cared about the human's goal. On average, participants provided an average rating for the emotion-aware robot that was 1.5 points higher than that for the emotion-ignorant robot. Question 10 measured the human perception of the robot's commitment to the collaboration; for this measure, the average participant score assigned to the emotion-aware robot was 6.2 points out of a maximum of 7 points, indicating that the participants felt that the emotion-aware robot was strongly committed to the collaboration. The emotion-ignorant robot received an average rating of 4.4 points, indicating only moderate commitment. These results strongly support Hypothesis 5, which posits that humans will feel that the emotion-aware robot will better understand their goals than the emotion-ignorant robot.

6.1.6 Human's Feeling about the Collaboration

Questions 11 and 12 were designed to gauge how the human participants felt about the partnership within the collaboration and the outcome of the collaboration. In the emotion-aware condition, participants scored Questions 11 and 12 as 6.1 and 6.3, respectively, indicating a strong sense of pursuing mutually agreed-upon




Figure 5: Results of the Likert scale surveys for 14 selected questions. The p-value for the difference between the means for each question is <0.001 , except for Q11, which is 0.008.

goals and very high satisfaction with the overall collaboration. It is worth noting that the participants scored Question 11 for the emotion-ignorant case at 5.3 points, leading to a 0.8-point difference between the two conditions; this is the smallest gap that occurs between the results for any question, and indicates that the participants felt that the goals were still partially mutually agreed-upon in the emotion-ignorant case. However, the general satisfaction with the outcome of the collaboration was significantly less in the case of the emotion-ignorant robot, at 4.8 points. These results support Hypothesis 6 that humans will feel greater a greater sense of mutual collaboration and understanding about the collaboration with the emotion-aware robot.

6.1.7 Human Perception of Mutual Satisfaction with Collaborative Partner

Questions 13 and 14 were designed to measure the human's perception of the robot's satisfaction with the human, and the human's satisfaction with the robot, respectively. The participants provided an average response in the emotion-aware condition of 5.8 and 5.9 to Questions 13 and 14, respectively, indicating a high level of mutual satisfaction; both answers were about 1.5 lower, on average, in the emotion-ignorant condition. These results indicate a higher level of satisfaction working with the robot in the emotion-aware condition, and strongly support Hypothesis 7, which posited that humans will feel a greater sense of mutual satisfaction with the emotion-aware robot than the emotion-ignorant robot.

6.2 Results from the Open-Ended Questionnaire

As described in Section 5.2, each participant answered an open-ended questionnaire at the end of the study. Figure 7 summarizes the questionnaire and which run users preferred for certain conditions (i.e. emotion-ignorant or emotion-aware). Note that some users either chose not to state a preference, or provided ambiguous answers regarding which run they preferred for certain conditions; these results were removed from the analysis. As shown in Figure 7, 100% of users unambiguously preferred the run with the emotion-aware robot. In general, this preference stemmed from a feeling of closeness and partnership, as seen in these responses: "the robot had emotions and responded to my emotions. Also, what

it said about my failing was cute and aimed to make me feel better." Another example is "I liked feeling needed and accounted for; I felt closer to the robot." Finally, "I saw the changes in its feeling, which motivated me to care more about my act...I also liked that he asked me to correct its failure, although it could ask the supervisor."

When asked in which of the two runs the robot exhibited behavior that could be useful in a more complex task, 93.75% chose the emotion-aware robot. In general, respondents thought that the emotion-aware robot was better at problem solving, more adaptable, and more capable of handling the social complexities that occur in collaboration, as shown in responses such as "The robot explained motives...which is important to keep a team communicating and on the same pace." Also, "When we failed he initially switched to a new task and then came back to the originally failed task. It kept me from getting irritated and negative." Finally, "The more complex, the more necessary it is to understand how humans think and operate...an empathetic robot can adapt, encourage and help." It is worth noting that one respondent preferred the emotion-ignorant case, saying "In a more complex task it might be better for the robot to take control and simply tell me what to do; trying to be understanding and collaborative wouldn't be as important as doing the task correctly."

The only question that did not provide statistically significant support in favor of the emotion-aware robot related to which case the robot exhibited behavior that could prevent human error. About 40% of respondents thought that the emotion-ignorant robot was more likely to prevent human error; however, all but one of these cited calling the supervisor as the main method of preventing human error, in spite of the fact that the instructions indicated that the robot's need to call the supervisor counted against the collaboration. Of the 60% who thought that the emotion-aware robot was better at preventing human error, most cited the robot's ability to console the human as the main behavior that could prevent human error. Respondents indicated that this enabled them to move on and feel better about the collaboration, as with this response: "The robot switched to a different task and we came back to an error later. This allowed my mind to move away from being frustrated. I was able to complete a different task which felt like a win - then come back and finish the error. Making my mind move away from frustration could definitely prevent more errors."

Question Category	Question	Question Number
Likability	I would like to continue working with the robot.	Q1
	I felt close to the robot.	Q2
Trust	I trust the robot.	Q3
	I trust the robot to perform appropriately in our collaboration.	Q4
Robot's Performance	The robot was repetitive.	Q5
	The robot's decisions improved my performance during the collaboration.	Q6
Robot's Understanding of Human's Emotions	The robot understood my emotions.	Q7
	The robot understands some of my feelings and takes them into account in our collaboration.	Q8
Robot's Understanding of Goals	The robot perceives accurately what my objectives are.	Q9
	The robot was committed to the collaboration.	Q10
Human Feeling about Collaboration	The robot and I are working towards mutually agreed-upon goals.	Q11
	I am satisfied with the outcome of our collaboration.	Q12
Satisfaction of Collaborative Partner	The robot was satisfied with my collaborative behavior.	Q13
	I was satisfied with the robot.	Q14

Figure 6: The 14 Likert scale questions presented in this paper, organized according to their groups.

Question	Percentage of Participants Favoring Emotion-Aware Robot	p-value
Which of the two runs with the robot did you prefer?	100%	0
In which of the two runs did the robot exhibit behavior that could be useful in a more complex task?	93.75%	$\ll 0.001$
In which of two runs did the robot exhibit behavior that could prevent human error?	60%	> 0.1
In which of the two runs did the robot exhibit behavior that could improve the efficiency of collaboration?	83.9%	$\ll 0.001$
What was the most interesting behavior of the robot and in which run did it happen?	82.8%	$\ll 0.001$

Figure 7: Open-ended questionnaire questions and results.

When asked in which of the runs the robot exhibited behavior that could improve the efficiency of the collaboration, 83.9% responded with the emotion-aware case; of these, the vast majority stated that this was because of the robot's ability to change the order of tasks in the event of a failure, and to ask the human for help.

Finally, when asked in which run the most interesting behavior occurred, 82.8% chose the emotion-aware condition. Of these respondents, 12 individuals stated that the robot's attempt to console the human by saying "It was not your fault" in response to the human's negative emotion that occurred as a consequence of the human's failed task was the most interesting behavior, and a majority mentioned that it actually made them feel more positive. Six participants referred to the robot's ability to understand and express emotion. Several participants referred to the robot's ability to communicate, including the ability to ask questions. Of those who responded with the emotion-ignorant case, most found the ability to call the supervisor, and mechanical functions, such as gripping, to be most interesting.

6.3 Impact of Demographics

As mentioned in Section 5.4, we recorded certain demographic information from each participant, including age and gender. We also had each participant complete several personality questionnaires. Although it was not the primary purpose of the study, we investigated the Likert scale results to determine if there were any

relevant trends based on the demographics and personalities of the participants. A close study of the results did not reveal any identifiable pattern based on gender or personality.

Age did reveal an interesting pattern. We divided the participants into two groups, below 30 years of age and 30 or above. While question-by-question comparisons revealed only a few statistically significant differences based on age, a consistent pattern emerged. For each of the fourteen questions presented, the younger age group reported higher scores for the emotion-aware robot (except Question 5, which is a reverse-score question). In the emotion-ignorant case, the younger group still scores the robot higher than the older group for 9 questions; for the other 5 questions, the older group scores the emotion-ignorant robot higher. In fact, a pattern emerged in which the score drop between the emotion-aware and the emotion-ignorant case was more for the younger group than for the older group; only questions 4, 6 and 12 broke this pattern.

7. DISCUSSION

Based on the results, all participants prefer to work with the emotion-aware robot. Humans find the emotion-aware robot more likable and more trustworthy, as indicated in the Likert-scale responses and the open-ended questionnaire responses. Based on the responses, the emotional interaction with the robot can help create a sense of closeness and enjoyment that makes humans want to continue working with the robot.




Figure 8: Impact of age on results of Likert scale Questions 1-7




Figure 9: Impact of age on results of Likert scale Questions 8-14

The results also indicate that the emotion aware robot can better maintain a collaborative relationship. Both Likert-scale responses, see Sections 6.1.5 and 6.1.7 and Open-Ended Questionnaire responses indicate this. Humans felt a stronger sense of the robot's commitment to the collaboration, and greater understanding of their goals and emotions from the robot. Several open-ended responses also indicated that the robot was able to successfully motivate people and maintain their commitment to the collaboration, especially when tasks failed. Additionally, as shown in Section 6.1.3, humans rated the emotion-aware case much higher than the emotion-ignorant case when asked which robot's decisions improved their performance, in essence acknowledging that their collaborator's (i.e. the robot's) decisions had a significant impact on their performance. As some of the open-ended responses indicated, successfully managing emotions within the collaboration can help keep the collaboration on track, and prevent distractions due to guilt and other negative emotions.

Finally, the emotion-aware robot developed a stronger sense of partnership through greater communication. The participants felt better understood by the emotion-aware robot, and felt that the goals were more mutually agreed-upon, refer to Sections 6.1.6 and 6.1.4, respectively. As evidenced in the following response, the emotion-aware robot was successfully able to create a sense of part-

nership through its more open communication style: "Communication is very important. In the first run (i.e. emotion-aware) the robot states what tasks he is working on, it is clear and straight-forward. Also during the first run the robot cares about the human(me)'s feelings and cheers me up when I failed at the tasks, I think that could also improve efficiency of collaboration, because it would be more like a team or partnership."

8. CONCLUSIONS

The goal of our user-study was to compare different aspects of humans' preferences during collaboration with a robot. Our results conclusively showed that humans prefer collaborating with an emotion-aware robot which is capable of a) expressing appropriate emotions, and b) changing collaborative behaviors based on the perceived negative emotions of the human. We are interested to carry out further studies with more capabilities from our framework in order to study collaboration with more complex tasks and evaluate collaborative performance based on objective measures such as cost or time to achieve a shared goal. We believe that emotion-aware robots will outperform emotion-ignorant robots in collaboration with humans.

9. REFERENCES

- [1] C. Adam and E. Lorini. A BDI emotional reasoning engine for an artificial companion. In *Workshop on Agents and multi-agent Systems for AAL and e-HEALTH (PAAMS)*, volume 430, pages 66–78. Springer, 2014.
- [2] C. Breazeal. *Designing Sociable Robots*. MIT Press, 2002.
- [3] C. Breazeal. Role of expressive behaviour for robots that learn from people. *Philosophical transactions of the Royal Society of London. Series B, Biological sciences*, 364(1535):3527–38, 2009.
- [4] C. Breazeal, A. Takanishi, and T. Kobayashi. *Social Robots that Interact with People*, pages 1349–1369. Springer Berlin Heidelberg, Berlin, Heidelberg, 2008.
- [5] D. Cameron, S. Fernando, E. Collins, A. Millings, R. Moore, A. Sharkey, V. Evers, and T. Prescott. Presence of life-like robot expressions influences children’s enjoyment of human-robot interactions in the field. In *4th International symposium on New Frontiers in Human-Robot Interaction.*, Canterbury, UK, 2015. M. Salem, A. Weiss, P. Baxter, & K. Dautenhahn (Eds.).
- [6] A. C. Gonzalez, M. Malfaz, and M. A. Salichs. An autonomous social robot in fear. *IEEE Transactions Autonomous Mental Development*, 5(2):135–151, 2013.
- [7] J. Gratch and S. C. Marsella. A domain-independent framework for modeling emotion. *Cognitive Systems Research*, 5(4):269–306, 2004.
- [8] B. J. Grosz and C. L. Sidner. Plans for discourse. In P. R. Cohen, J. Morgan, and M. E. Pollack, editors, *Intentions in Communication*, pages 417–444. MIT Press, Cambridge, MA, 1990.
- [9] A. I. Guha and S. Tellex. Towards Meaningful Human-Robot Collaboration on Object Placement. In *RSS Workshop on Planning for Human-Robot Interaction: Shared Autonomy and Collaborative Robotics*, 2016.
- [10] H.-R. Kim and D.-S. Kwon. Computational model of emotion generation for human-robot interaction based on the cognitive appraisal theory. *Journal of Intelligent and Robotic Systems*, 60(2):263–283, 2010.
- [11] M. Klug and A. Zell. Emotion-based human-robot-interaction. In *IEEE 9th International Conference on Computational Cybernetics (ICCC)*, 2013.
- [12] I. Leite, A. Pereira, S. Mascarenhas, C. Martinho, R. Prada, and A. Paiva. The influence of empathy in human-robot relations. *International Journal of Human-Computer Studies*, 71(3):250–260, 2013.
- [13] A. Lim and H. G. Okuno. The mei robot: Towards using motherese to develop multimodal emotional intelligence. *IEEE Transactions Autonomous Mental Development*, 6(2):126–138, 2014.
- [14] C. Liu and N. Sarkar. Online affect detection and robot behavior adaptation for intervention of children with autism. *IEEE TRANSACTIONS ON ROBOTICS*, 24(4):883–896, 2008.
- [15] R. Looije, M. Neerincx, and G.-J. Kruijff. Affective collaborative robots for safety crisis management in the field. In *Proceedings of the 4th International Conference on Information Systems for Crisis Response and Management (ISCRAM 2007)*. ISCRAM, 5 2007.
- [16] R. P. Marinier III, J. E. Laird, and R. L. Lewis. A computational unification of cognitive behavior and emotion. *Cognitive System Research*, 10(1):48–69, March 2009.
- [17] S. C. Marsella and J. Gratch. EMA: A process model of appraisal dynamics. *Cognitive Systems Research*, 10(1):70–90, March 2009.
- [18] A. Paiva, I. Leite, and T. Ribeiro. Emotion modeling for sociable robots. In J. G. A. K. Rafael A. Calvo, Sidney D’Mello, editor, *Handbook of Affective Computing*, pages 296–308. Oxford University Press, 2014.
- [19] M. Pontier and J. F. Hoorn. How women think robots perceive them - as if robots were men. In *International Conference on Agents and Artificial Intelligence (ICAART-2)*, pages 496–504, 2013.
- [20] M. Shayganfar, C. Rich, and C. L. Sidner. A design methodology for expressing emotion on robot faces. In *IROS*, pages 4577–4583. IEEE, 2012.
- [21] M. Si, S. C. Marsella, and D. V. Pynadath. Modeling appraisal in theory of mind reasoning. *Autonomous Agents and Multi-Agent Systems*, 20(1):14–31, 2010.
- [22] Suppressed for Anonymity. 2015.
- [23] Suppressed for Anonymity. 2016.
- [24] Suppressed for Anonymity. 2016.
- [25] Suppressed for Anonymity. 2016.
- [26] D. Vogiatzis, C. Spyropoulos, V. Karkaletsis, Z. Kasap, C. Matheson, and O. Deroo. An affective robot guide to museums. In *Proceedings of the 4th International Workshop on Human-Computer Conversation*, 2008.
- [27] T. Wehrle. Motivations behind modeling emotional agents: Whose emotion does your robot have?, 1998.
- [28] T. Zhang, D. B. Kaber, B. Zhu, M. Swangnetr, P. Mosaly, and L. Hodge. Service robot feature design effects on user perceptions and emotional responses. *Intelligent Service Robotics*, 3(2):73–88, 2010.